

Class 7:Machine Learning 1

Yerin Go (A16272901)

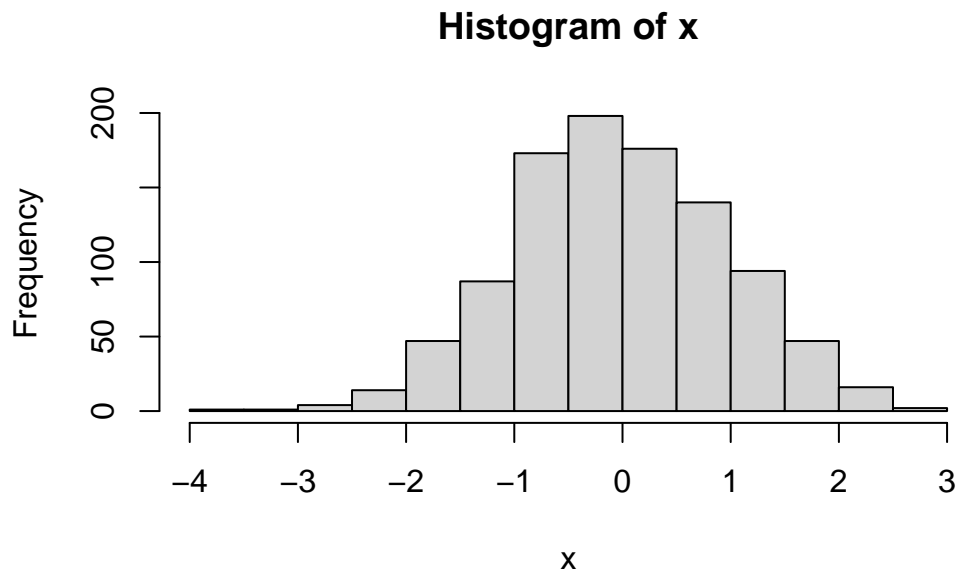
#Clustering Methods

The broad goal here is to find groupings (clusters) in your input data

##kmeans

First, let's make up some data to cluster.

```
x <- rnorm(1000)
hist(x)
```



Make a vector of length 60 with 30 points centered at -3 and 30 points centered at +3

```
tmp <- c(rnorm(30, mean = -3), rnorm(30, mean = 3))
tmp
```

```
[1] -3.552130 -3.986926 -4.352605 -1.134901 -1.789953 -2.340931 -4.161691
[8] -2.855190 -1.470861 -2.659238 -3.514743 -4.112074 -2.842861 -3.536923
[15] -4.008434 -3.952195 -4.196491 -3.601648 -3.950042 -1.098229 -4.237116
[22] -2.539294 -4.300440 -4.288732 -1.924482 -3.687145 -3.019116 -5.292654
[29] -2.622738 -1.878402  3.302369  1.615299  2.204554  5.286726  2.748953
[36]  4.012986  3.015273  3.140110  5.248738  2.862859  4.136961  3.848212
[43]  1.641793  2.797816  5.416417  2.419450  1.945332  2.819614  3.889826
[50]  3.401286  3.181459  4.889458  2.482277  1.519214  2.436674  2.262992
[57]  2.167018  2.982171  4.286173  4.357386
```

I will not make a wee x and y dataset with 2 groups of points.

```
rev(c(1:5))
```

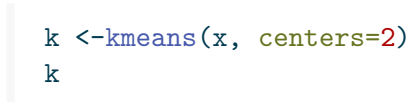
```
[1] 5 4 3 2 1
```

```
x <- cbind(x=tmp, y=rev(tmp))
x
```

```
      x      y
[1,] -3.552130  4.357386
[2,] -3.986926  4.286173
[3,] -4.352605  2.982171
[4,] -1.134901  2.167018
[5,] -1.789953  2.262992
[6,] -2.340931  2.436674
[7,] -4.161691  1.519214
[8,] -2.855190  2.482277
[9,] -1.470861  4.889458
[10,] -2.659238  3.181459
[11,] -3.514743  3.401286
[12,] -4.112074  3.889826
[13,] -2.842861  2.819614
[14,] -3.536923  1.945332
[15,] -4.008434  2.419450
[16,] -3.952195  5.416417
```

[17,]	-4.196491	2.797816
[18,]	-3.601648	1.641793
[19,]	-3.950042	3.848212
[20,]	-1.098229	4.136961
[21,]	-4.237116	2.862859
[22,]	-2.539294	5.248738
[23,]	-4.300440	3.140110
[24,]	-4.288732	3.015273
[25,]	-1.924482	4.012986
[26,]	-3.687145	2.748953
[27,]	-3.019116	5.286726
[28,]	-5.292654	2.204554
[29,]	-2.622738	1.615299
[30,]	-1.878402	3.302369
[31,]	3.302369	-1.878402
[32,]	1.615299	-2.622738
[33,]	2.204554	-5.292654
[34,]	5.286726	-3.019116
[35,]	2.748953	-3.687145
[36,]	4.012986	-1.924482
[37,]	3.015273	-4.288732
[38,]	3.140110	-4.300440
[39,]	5.248738	-2.539294
[40,]	2.862859	-4.237116
[41,]	4.136961	-1.098229
[42,]	3.848212	-3.950042
[43,]	1.641793	-3.601648
[44,]	2.797816	-4.196491
[45,]	5.416417	-3.952195
[46,]	2.419450	-4.008434
[47,]	1.945332	-3.536923
[48,]	2.819614	-2.842861
[49,]	3.889826	-4.112074
[50,]	3.401286	-3.514743
[51,]	3.181459	-2.659238
[52,]	4.889458	-1.470861
[53,]	2.482277	-2.855190
[54,]	1.519214	-4.161691
[55,]	2.436674	-2.340931
[56,]	2.262992	-1.789953
[57,]	2.167018	-1.134901
[58,]	2.982171	-4.352605
[59,]	4.286173	-3.986926

```
plot(x)
```



Cluster means:

Clustering vector:

Within cluster sum of squares by cluster:

4

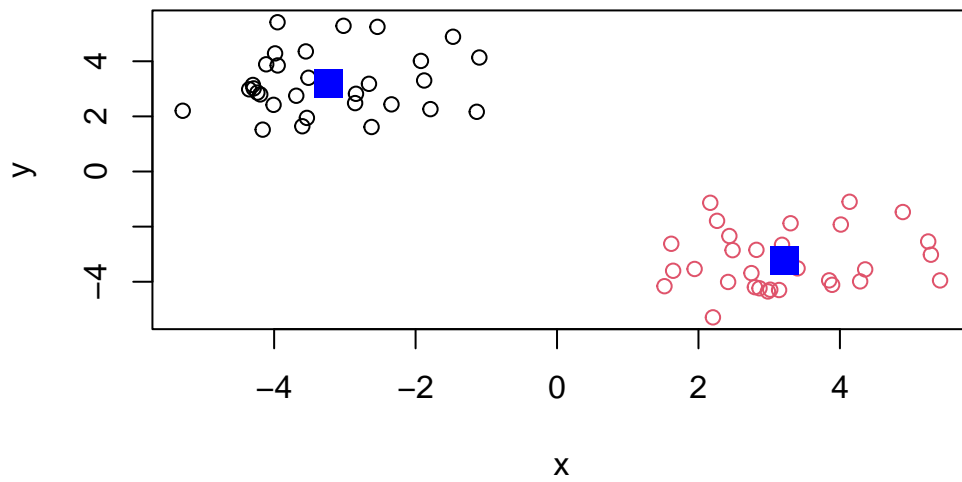
```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

k\$size

```
k$cluster
```

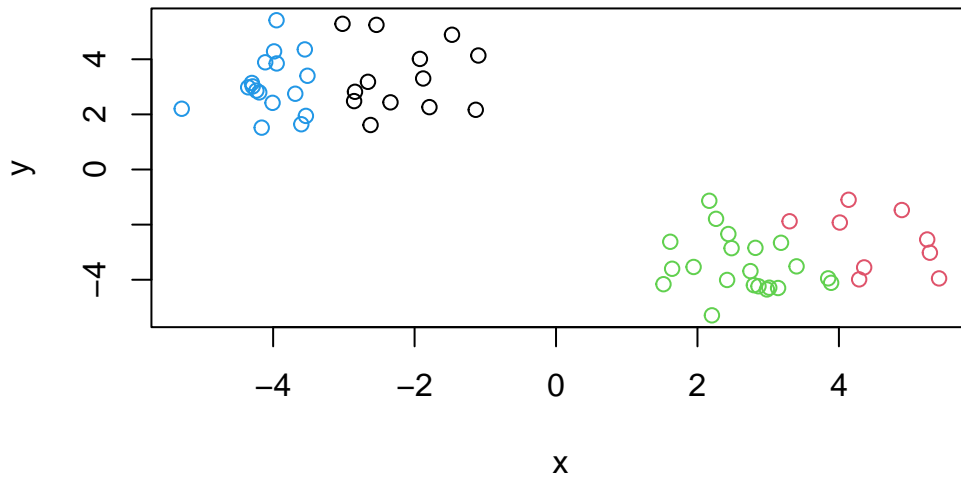
k\$centers

```
plot(x, col=k$cluster)
points(k$centers, col="blue", pch= 15, cex=2)
```



We can cluster into 4 groups.

```
# kmeans
k4 <- kmeans(x, centers=4)
#plot results
plot(x, col=k4$cluster)
```



A big limitation of `kmeans` is that it does what you ask even if you ask for silly clusters. # Hierarchical Clustering

The main base R function for Hierarchical Clustering is `hclust()` Unlike `kmeans()` you cannot just pass it your data as input. You first need to calculate a distance matrix.

```
d <- dist(x)
hc <- hclust(d)
hc
```

Call:

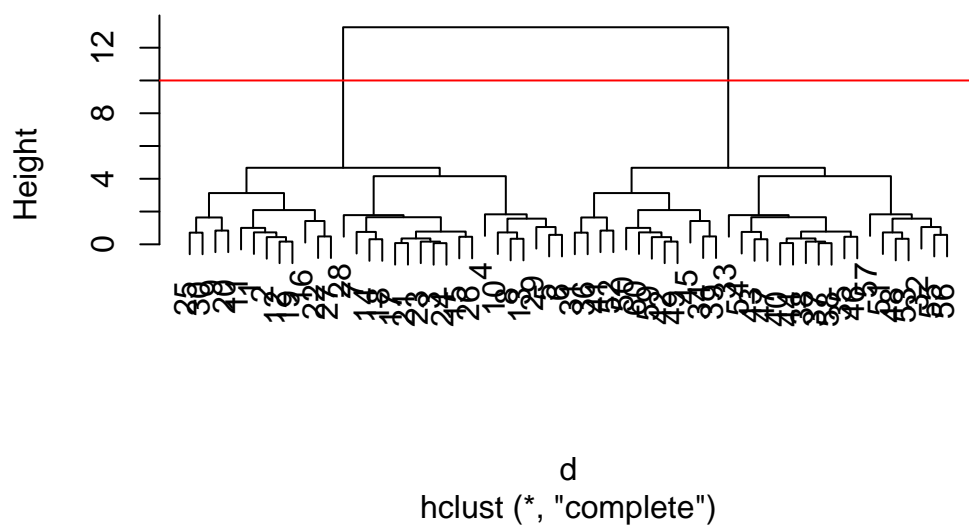
```
hclust(d = d)
```

```
Cluster method : complete
Distance       : euclidean
Number of objects: 60
```

Use `plot()` to view results.

```
plot(hc)
abline(h=10, col="red")
```

Cluster Dendrogram



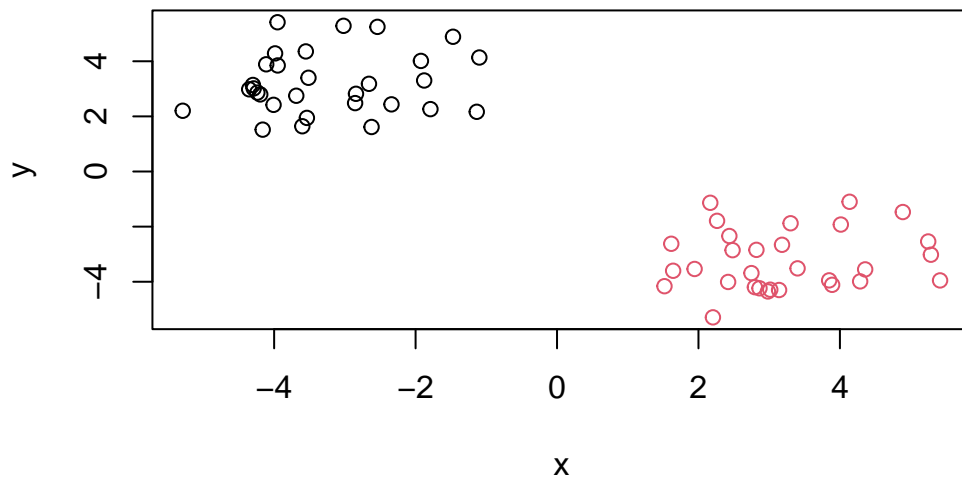
To make the “cut” and get our cluster membership vector we can use the `cutree()` function.

```
grps <- cutree(hc, h=10)
grps
```

[illegible]

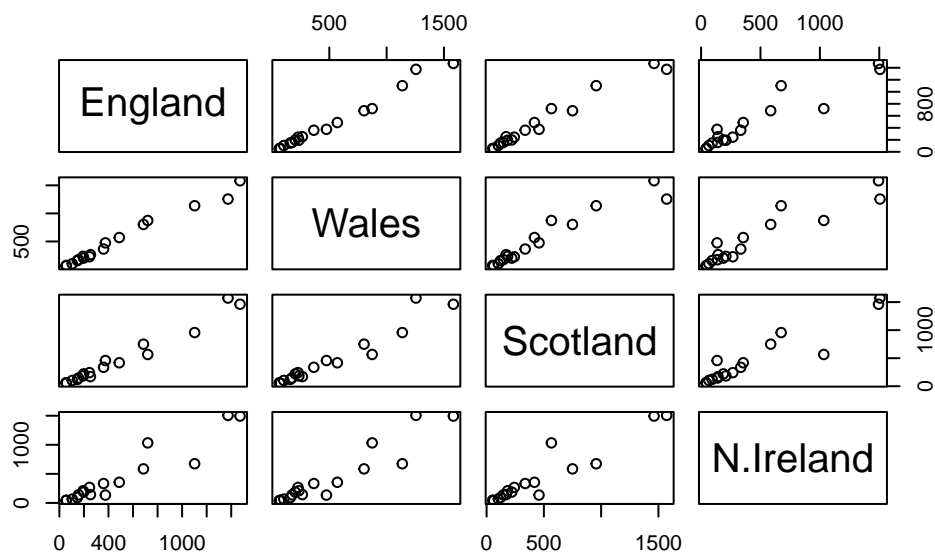
Make a plot of our data colored by hclust results.

```
plot(x, col=grps)
```

#Principle Component Analysis (PCA) Here we will do PCA on some food data from the UK.

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names=1)
plot(x)
```



#Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
dim(x)
```

```
[1] 17  4
```

```
nrow(x)
```

```
[1] 17
```

```
ncol(x)
```

```
[1] 4
```

##PCA to the rescue The main “base” R function for PCA is called `prcomp()` Here we need to take the transpose of our input as we want the countries in the rows and foods as the columns.

```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

. Q. How much variance is captured to 2 PCs 96.5%

To make our main “PC score plot” or “PC1 vs. PC2 plot” or “PC plot” or “Ordination Plot”.

```
attributes(pca)
```

\$names

```
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

\$class

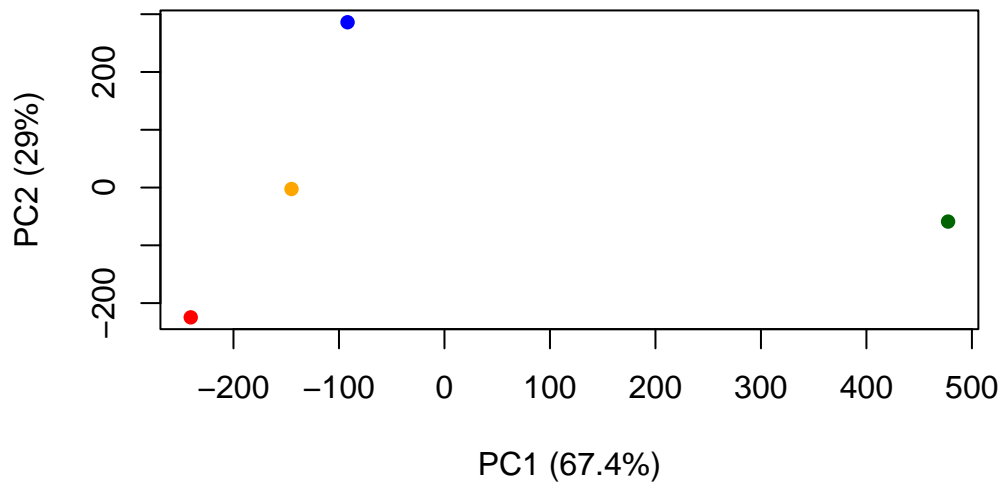
```
[1] "prcomp"
```

We are after the `pca$x` result component to make out main PCA plot

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-9.152022e-15
Wales	-240.52915	-224.646925	-56.475555	5.560040e-13
Scotland	-91.86934	286.081786	-44.415495	-6.638419e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.329771e-13

```
mycols <- c("orange", "red", "blue", "darkgreen")
plot(pca$x[,1], pca$x[,2], col=mycols, pch=16, xlab= "PC1 (67.4%)", ylab= "PC2 (29%)")
```



Another important result from PCA is how the original variables (in this case the foods) contribute to the PCs. This is contained in the `pca$rotation` object- folks often call this the “loadings” or “contributions” to the PCs.

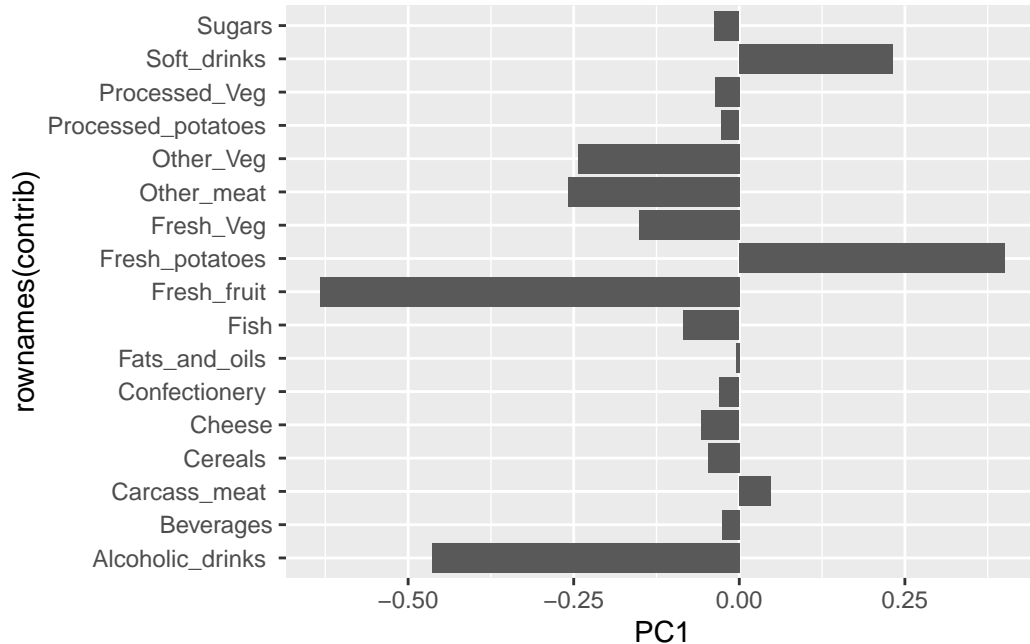
```
pca$rotation[,1]
```

Cheese	Carcass_meat	Other_meat	Fish
-0.056955380	0.047927628	-0.258916658	-0.084414983
Fats_and_oils	Sugars	Fresh_potatoes	Fresh_Veg
-0.005193623	-0.037620983	0.401402060	-0.151849942
Other_Veg	Processed_potatoes	Processed_Veg	Fresh_fruit
-0.243593729	-0.026886233	-0.036488269	-0.632640898
Cereals	Beverages	Soft_drinks	Alcoholic_drinks
-0.047702858	-0.026187756	0.232244140	-0.463968168
Confectionery			
-0.029650201			

We can make a plot along PC1.

```
library(ggplot2)
contrib <- as.data.frame(pca$rotation)
```

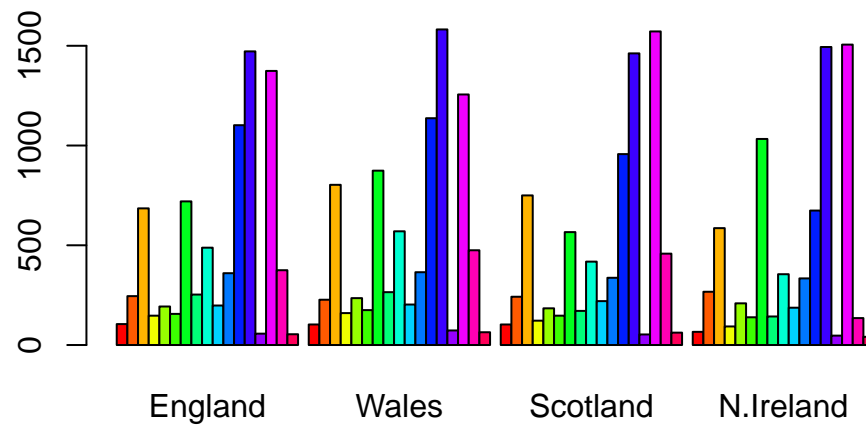
```
ggplot(contrib) +
  aes(PC1, rownames(contrib))+
  geom_col()
```



#Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

I prefer adding the part of the code `row.names=1` to the of the line of code `x <- read.csv(url, row.names=1)`. This approach is more robust than using the `rownames(x) <- x[,1]` because everything I run the code the second way, it will override the function, therefore deleting the relative first row each time the code is run.

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



#Q3: Changing what optional argument in the above `barplot()` function results in the following plot?

```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```

